Detailed Analysis of Advanced Bin Packing Algorithms: From Rothvoß (2013) to Hoberg & Rothvoß (2015)

Shreyas Ramasubramanian

November 5, 2025

Abstract

For over three decades, the $OPT + O(\log^2 OPT)$ guarantee of the Karmarkar-Karp algorithm (1982) stood as the benchmark for additive approximations to the 1D Bin Packing Problem. This review provides a detailed analysis of the two sequential breakthroughs that finally surpassed this bound. We examine the 2013 paper by Rothvoß, focusing on its introduction of discrepancy theory (via the Lovett-Meka algorithm) and the novel "gluing" technique used to handle problematic "spiky" patterns, achieving an $OPT + O(\log OPT \cdot \log\log OPT)$ bound. We then delve into the 2015 paper by Hoberg & Rothvoß, detailing its refined "cleaner" 2-stage packing mechanism (Items $\to$ Containers $\to$ Bins), explaining why this structural change simplifies the analysis and secures the final, tight $OPT + O(\log OPT)$ guarantee.

1 Introduction: The Quest for Additive Guarantees

The one-dimensional Bin Packing Problem (1DBPP) asks for the minimum number of unit-capacity bins required to pack a given set of items with sizes $s_1, \ldots, s_n \in (0, 1]$. Its NP-hardness necessitates approximation algorithms. While simple heuristics like First-Fit Decreasing (FFD) offer good *multiplicative* guarantees ($\approx 1.222 \cdot OPT + \text{const}$), a significant line of theoretical research seeks *additive* guarantees of the form $OPT + C$, where C is a small, slowly growing function, ideally independent of the number of items n.

For over 30 years, the landmark result was the **Karmarkar-Karp (KK82) algorithm**, achieving $OPT_f + O(\log^2 OPT_f)$ bins, where OPT_f is the optimal value of the fractional relaxation. This review focuses on the two papers that finally improved this bound.

2 Preliminaries: The Gilmore-Gomory LP Relaxation

Both KK82 and the subsequent breakthroughs leverage the **Gilmore-Gomory Linear Program (LP) relaxation**.

- **Structure:** The LP operates over *patterns*. A pattern p is a multiset of items that fits into a single bin, represented as a vector $p \in \mathbb{Z}^n_{\geq 0}$ where p_i is the count of item i, such that $\sum_{i=1}^n p_i s_i \leq 1$.
- Let $\mathcal{P}$ be the set of all possible valid patterns. The LP formulation is:

$$\min \sum_{p \in \mathcal{P}} x_p \quad \text{subject to} \quad \sum_{p \in \mathcal{P}} p_i x_p \geq b_i \quad \forall i = 1, \ldots, n, \quad x_p \geq 0 \quad \forall p \in \mathcal{P}$$

  Here, b_i is the required number of items of type i, and x_p is a variable representing how many times pattern p is used.

- **Challenge:** The number of possible patterns $|\mathcal{P}|$ can be exponential in n. However, the LP has only n constraints. Efficient algorithms (like Ellipsoid method variants or Column Generation, referenced in KK82 and Rothvoß's papers) can find an approximate solution x with cost $\leq OPT_f + \delta$ in polynomial time (polynomial in n, $\sum b_i$, and $1/\delta$). Often, one works with a basic feasible solution, which has at most n non-zero x_p values.

- **The Core Problem:** The LP yields a fractional optimum OPT_f. The challenge is to round the fractional vector x (with potentially millions of components x_p) into an integer vector y (representing an actual packing) such that the cost $\sum y_p$ is very close to $\sum x_p$, specifically $\sum y_p \leq \sum x_p + C$.

3 The 2013 Breakthrough: Rothvoß

Rothvoß's 2013 paper, "Approximating Bin Packing within $O(\log OPT \cdot \log\log OPT)$ bins," was the first improvement over KK82. It introduced discrepancy theory as a new rounding tool.

3.1 The New Tool: Discrepancy Theory via Lovett-Meka (LM12)

Instead of KK82's rounding, Rothvoß employed the **Constructive Partial Coloring Lemma** (Lovett & Meka, 2012).

- **Intuition:** Discrepancy theory aims to find a coloring (e.g., ± 1) for elements of a ground set such that for any given subset, the sum of colors is small (the set is "balanced"). The LM12 algorithm provides a constructive, randomized way to achieve this, generalized to rounding fractional vectors.
- **The LM12 Guarantee (Lemma 1 in Rothvoß 2013):** Given a starting fractional point $x \in [0,1]^m$ and constraints defined by vectors $v_1, \ldots, v_n$, the algorithm finds a new point $y \in [0,1]^m$ such that:
  1. **Near-Integrality:** At least half of the coordinates y_j are very close to 0 or 1 ($y_j \in [0,\delta] \cup [1-\delta, 1]$).
  2. **Low Error (Discrepancy):** The change in each constraint value is bounded: $|v_i y - v_i x| \leq \lambda_i ||v_i||_2$. The λ_i are parameters chosen based on an "entropy condition" ($\sum e^{-\lambda_i^2/16} \leq m/16$), ensuring the total allowed error isn't too large.
- **Iterative Rounding:** The algorithm applies LM12 repeatedly ($O(\log m)$ times) to eventually round all variables to 0 or 1.

Here is the formal statement of the lemma, which you can add for mathematical rigor:

Lemma 1 (Constructive Partial Coloring, Lovett-Meka 2012)**.** *Let $A \in \mathbb{R}^{n \times m}$ be a matrix, $x \in [0,1]^m$ be a fractional solution, and $\lambda_1, \ldots, \lambda_n > 0$ be parameters satisfying $\sum_{i=1}^n e^{-\lambda_i^2/16} \leq \frac{m}{16}$.*

There exists a polynomial-time randomized algorithm that finds a point $y \in [0,1]^m$ such that:

1. ***(Low Discrepancy):*** *The error on each constraint is bounded by the L_2-norm of the corresponding row:*
$$|(Ay)_i - (Ax)_i| \leq \lambda_i \|A_i\|_2 \quad \forall i = 1, \ldots, n$$
2. ***(Near-Integrality):*** *For a given $\delta > 0$, at least half of the coordinates are nearly integral. For $m' \geq (1-\delta)m$ coordinates $j \in [m]$, we have:*
$$y_j \in [0,\delta] \cup [1-\delta, 1]$$

Rothvoß (2013) uses this lemma iteratively. In each of the $O(\log m)$ iterations, half of the remaining fractional variables are rounded to 0 or 1, while the cumulative error on each constraint i is carefully controlled by the $\lambda_i \|A_i\|_2$ term.

3.2 The Core Challenge: "Spiky" Patterns and the L_2-Norm

The effectiveness of LM12 hinges on the error bound $|v_i y - v_i x| \leq \lambda_i ||v_i||_2$. Rothvoß's insight was that this bound is problematic if $||v_i||_2$ is large.

- In the bin packing context, the vectors v_i correspond to (sums of) rows of the pattern matrix A. The i-th row A_i lists how many times item i appears in each pattern.
- A large $||A_i||_2 = \sqrt{\sum_p (A_{ip})^2}$ norm occurs if there are patterns p with large entries A_{ip} — i.e., patterns containing many copies of item i.
- This is the **"spiky pattern"** problem: a pattern p heavily utilizing a single item type i (large A_{ip}) causes a large L_2-norm for row i, leading to a potentially large rounding error for that item's constraint. This was particularly problematic for very small items where A_{ip} could be large.

3.3 The Novel Technique: "Gluing" (Section 5.2)

To mitigate the large L_2-norms caused by spiky patterns, Rothvoß introduced **"gluing"** as a pre-processing step before rounding.

- **Trigger:** Gluing is applied when a pattern p (with fractional value $x_p = r/q$) uses many copies (p_i) of a *small* item i, such that the total size $p_i s_i$ exceeds a threshold (e.g., related to $1/\text{polylog}(n)$). Specifically, if $p_i \geq w \cdot q$ for carefully chosen w, q.
- **Mechanism (Figure 1b):** It takes $w \cdot q$ copies of item i within pattern p and conceptually "glues" them into q copies of a *new, artificial, larger item* i' with size $s_{i'} = w \cdot s_i$. The pattern p is modified to use q copies of i' instead. The fractional value x_p remains r/q.
- **Purpose:** This transformation replaces a large entry A_{ip} (for the small item i) with smaller entries (q) for the new, larger item i'. This "smooths" the matrix rows associated with small items, reducing their L_2-norms and making the LM12 rounding effective.
- **Limitation:** As noted by Hoberg & Rothvoß (2015), this gluing procedure in the 2013 paper was complex and primarily effective only for items below a size threshold of $1/\text{polylog}(n)$.

3.4 The Full Algorithm (Section 6) and Result

The Rothvoß (2013) algorithm iterates $O(\log n)$ times. Each iteration involves:

1. **Discretize (Lemma 10):** Ensure x_p values are multiples of some $1/q = 1/\text{polylog}(n)$.
2. **Group & Glue (Lemma 8):** Apply standard grouping and the novel gluing to make the instance "well-spread" for small items (i.e., ensure A_{ip} is small relative to the total count $A_i x$ for small i).
3. **Round (Theorem 9):** Apply the LM12 algorithm to the modified matrix/solution, rounding half the remaining fractional variables with controlled error. The error analysis in Theorem 9 carefully bounds the discrepancy based on the properties achieved by gluing.

The complexity arose because the "Group & Glue" step introduced an error term related to the parameters used (specifically, involving $\log(1/\delta)$ where δ related to $1/\text{polylog}(n)$), contributing an $O(\log \log OPT)$ factor to the error accumulated in each of the $O(\log OPT)$ iterations. The final guarantee: $OPT_f + O(\log OPT \cdot \log \log OPT)$. The runtime is polynomial, dominated by solving the LP and the LM12 calls.

4 The 2015 Refinement: Hoberg & Rothvoß

The 2015 Hoberg & Rothvoß paper, "A Logarithmic Additive Integrality Gap for Bin Packing," achieved the tighter $O(\log OPT)$ bound by introducing a fundamentally cleaner structure.

4.1 The Core Idea: 2-Stage Packing (Section 2)

Instead of fixing the spiky pattern problem after the fact, they reformulated the problem to avoid it structurally.

1. **Stage 1: Items → Containers.** Define a **container** C as any valid multiset of original items that fits in a bin ($C \in \mathbb{Z}^n_{\geq 0}$ with $\sum s_i C_i \leq 1$). Let $s(C)$ be its total size.

2. **Stage 2: Containers → Bins.** Define a **pattern** p now as a multiset of *containers* that fits in a bin ($p \in \mathbb{Z}^{\mathcal{C}}_{\geq 0}$ where $\mathcal{C}$ is the set of all containers, such that $\sum_{C \in \mathcal{C}} p_C s(C) \leq 1$).

3. **Intermediate Variables:** They introduce integer variables y_C representing the number of times container C is "created" or used.

4. **Packing Graphs (Figure 1):** The process is modeled with two bipartite graphs:
   - $G_1(b, y)$: Matches original items (demand b_i) to available slots within the chosen containers (supply $y_C \cdot C_i$).
   - $G_2(x, y)$: Matches the chosen containers (demand y_C) to slots within the final fractional patterns (supply $x_p \cdot p_C$).

5. **Deficiency:** The objective is implicitly tied to minimizing a "deficiency" metric, $\text{def}(x, y) = \text{def}(G_1) + \text{def}(G_2)$, which measures the total size of items/containers left unpacked by the fractional assignments in G_1 and G_2. The goal of the rounding is to keep the final deficiency small.

This 2-stage packing concept is formalized into a new, structured LP relaxation. Let $\mathcal{C}$ be the set of all valid "containers" (multisets of items that fit in a bin), and let $\mathcal{P}_{\mathcal{C}}$ be the set of all valid "patterns" of containers (multisets of containers that fit in a bin).

The Hoberg & Rothvoß (2015) algorithm finds a fractional solution (x, y) to the following system:

The 2-Stage Packing Formulation

$$\min \sum_{p \in \mathcal{P}_{\mathcal{C}}} x_p \tag{1}$$

$$\text{s.t.} \sum_{C \in \mathcal{C}} y_C \cdot C_i \geq b_i \qquad \forall \text{ items } i \tag{2}$$

$$\sum_{p \in \mathcal{P}_{\mathcal{C}}} x_p \cdot p_C \geq y_C \qquad \forall \text{ containers } C \in \mathcal{C} \tag{3}$$

$$x_p \geq 0, \quad y_C \geq 0 \qquad \forall p \in \mathcal{P}_{\mathcal{C}}, C \in \mathcal{C}$$

Here, y_C represents the (fractional) number of times container C is "created", and x_p is the number of times pattern p is used.

- Equation (2) ensures that all original items b_i are packed into containers (corresponds to graph G_1).
- Equation (3) ensures that all created containers y_C are packed into the final bins (corresponds to graph G_2).

The key insight is that the rounding algorithm is only applied to Equation (3), which packs containers. As noted, a pattern p simply cannot contain an arbitrarily large number of copies (p_C) of the same container C. This structurally avoids the "spiky" L_2-norm problem, as the matrix for (3) is inherently "smooth".

4.2 The Final Algorithm and Bound

The Hoberg & Rothvoß (2015) algorithm also iterates $O(\log n)$ times:

1. **Rebuild Containers (Section 3):** This step replaces the complex "gluing". It involves sophisticated grouping (Lemma 10) and reassigning/combining containers (Lemma 13) within the existing fractional patterns x. The goal is to ensure the container patterns have nice structural properties (e.g., bounding the number of times a single container type appears in one pattern, $p_C \leq (1/\sigma)^{1/4}$ for containers in size class σ). This makes the matrix A (now representing container incidences) well-behaved for rounding.

2. **Round (Section 4):** Apply the *same* LM12 algorithm (Claim 14) to the fractional pattern vector x, using constraints derived from the (rebuilt) container incidence matrix A.

4.3 Why it's Cleaner and Better

The 2-stage structure provides several advantages leading to the cleaner analysis and tighter bound:

- **Inherent Smoothness:** The objects being packed into the final patterns are now "containers." Since containers must have size $s(C) \leq 1$, a pattern p simply cannot contain an arbitrarily large number of copies (p_C) of the same container C. The large entries in the matrix rows that plagued the 2013 analysis for tiny items are structurally avoided when dealing with containers.

- **Simplified Pre-processing:** The "Rebuilding Containers" step (Section 3) achieves the necessary "well-spread" properties more elegantly and robustly than the 2013 "gluing," which had limitations on item size. The 2015 paper notes their procedure works even for items/containers up to size $\Omega(1)$.

- **Error Reduction ($O(1)$ Deficiency):** Because the matrix A (representing container patterns) is inherently better behaved, the LM12 rounding step (Section 4) incurs only a constant $O(1)$ increase in total deficiency per iteration (Lemma 17 and subsequent analysis). The complex error term involving $\log\log OPT$ from the 2013 analysis disappears.

- **Full Spectrum Parameters:** The 2015 paper notes they use the "full spectrum" of error parameters λ_I in LM12, whereas the 2013 paper used only two types, contributing to the cleaner $O(1)$ error per iteration.

The final result: The total additive gap is $O(\log OPT)$ iterations $\times$ $O(1)$ error per iteration, yielding the tight $\mathbf{OPT_f} + \mathbf{O}(\log \mathbf{OPT})$ bound.

5 Conclusion

The progression from Rothvoß (2013) to Hoberg & Rothvoß (2015) exemplifies refinement in theoretical algorithm design. Both papers leverage the Gilmore-Gomory LP and the powerful Lovett-Meka discrepancy rounding algorithm. However, the 2013 paper required a complex, somewhat limited "gluing" mechanism to force the problem structure into shape for the rounding tool, resulting in an $O(\log OPT \cdot \log\log OPT)$ gap. The 2015 paper achieved the final $O(\log OPT)$ gap through a more fundamental insight: redefining the problem via a 2-stage packing (Items $\rightarrow$ Containers $\rightarrow$ Bins). This elegant structural change inherently created the "smoothness" needed for discrepancy rounding, leading to a cleaner algorithm, a simpler analysis, and a tighter, likely optimal, additive bound. Visualizing the contrast between "gluing" and "2-stage packing" will be key to understanding this significant theoretical advancement.

Formally, the two breakthrough results are:

Theorem 1 (Rothvoß 2013)**.** *There exists a polynomial-time algorithm that, given a bin packing instance I, finds a packing using at most*

$$OPT_f(I) + O(\log(OPT_f(I)) \cdot \log\log(OPT_f(I)))$$

bins, where $OPT_f(I)$ is the optimal value of the Gilmore-Gomory LP relaxation.

Theorem 2 (Hoberg & Rothvoß 2015)**.** *There exists a polynomial-time algorithm that, given a bin packing instance I, finds a packing using at most*

$$OPT_f(I) + O(\log(OPT_f(I)))$$

bins.